
Text Classification with Deep Learning: Character-Level Recurrent Networks vs. Parameter-Efficient Transformers

Anrui Wang
2023533015

wangar2023@shanghaitech.edu.cn

Abstract

We study binary sentiment classification on a Chinese review corpus, comparing a character-level bidirectional LSTM with a pretrained RoBERTa-based classifier. For the advanced task we extend the RoBERTa baseline with layer mixing, gated multi-view pooling, and parameter-efficient adaptation via LoRA and BitFit. Under a controlled matched-baseline protocol the advanced architecture improves validation macro F1 from 0.8622 to 0.9317 (+0.0695), reduces validation loss from 0.2166 to 0.1435, and achieves 95.80% accuracy, while training only 3.57M parameters out of 105.7M total—confirming that careful feature aggregation and targeted adaptation are more effective than naïve full fine-tuning.

1 Introduction

Sentiment analysis over Chinese-language text presents a practical and well-scoped setting for evaluating neural text classifiers. Unlike English, Chinese lacks whitespace-delimited tokens, which makes the choice of segmentation granularity—character, word, or subword—a consequential design decision. Moreover, the prevalence of short, colloquial reviews combined with a moderate degree of class imbalance demands that classifiers generalise effectively beyond the dominant class.

This project addresses binary sentiment classification on a Chinese review dataset and pursues two complementary goals. The *basic task* compares two substantially different architectural families: (i) a lightweight character-level bidirectional LSTM trained entirely from scratch, and (ii) a pretrained Chinese RoBERTa encoder fine-tuned with a simple pooled classification head. The *advanced task* then develops a stronger system by modifying the RoBERTa-based model through a combination of multi-view sentence-representation learning and parameter-efficient adaptation.

The main contributions of the report are as follows. First, we provide a rigorous empirical comparison of recurrent and Transformer architectures on the same dataset under identical evaluation conditions, demonstrating that pretrained contextual representations provide a substantial advantage. Second, we introduce and evaluate a controlled architectural modification that adds layer mixing, gated multi-view pooling, and LoRA-based adaptation on top of a frozen RoBERTa encoder, showing a macro F1 improvement of +0.0695 over a matched frozen-head baseline. Third, we show that the improvement survives migration from a hyperparameter-search sandbox into the project training pipeline, confirming reproducibility. Finally, we present a detailed analysis connecting the observed performance gains to specific architectural choices, informing future design decisions.

2 Dataset and Preprocessing

2.1 Dataset

The dataset consists of Chinese consumer reviews labelled for binary sentiment. Each example contains a single `review` text field and a binary `label` field. The training set contains 10,000 examples; the validation set contains 2,500. Table 1 summarises the key statistics.

Table 1: Dataset statistics.

Split	Samples	Positive	Negative	Pos. Ratio	Avg. Chars
Train	10,000	8,066	1,934	80.66%	53.29
Validation	2,500	2,025	475	81.00%	54.69

Two properties of the data drive the experimental design. First, the class distribution is moderately imbalanced: positive examples account for approximately 81% of both splits. A classifier that predicts the majority class always would achieve ~81% accuracy yet zero minority-class recall, making accuracy an insufficient primary metric. We therefore adopt *macro F1* as the principal model-selection criterion throughout. Second, the average review length is only 53–55 characters, but a small number of examples exceed 1,000 characters, necessitating explicit sequence truncation.

2.2 Preprocessing Pipeline

Character-level preprocessing (BiLSTM). The LSTM model operates directly on the raw character sequence. A character vocabulary is constructed from the training set with a minimum frequency threshold of 2; characters appearing fewer than twice are mapped to a shared `<unk>` token. The resulting vocabulary size is 3,027 including the `<pad>` and `<unk>` special tokens. All sequences are truncated or padded to a maximum length of 512 characters. No lowercasing or normalisation beyond `<unk>` substitution is applied.

Subword preprocessing (BERT models). Both RoBERTa-based models use the Hugging Face tokenizer shipped with `hfl/chinese-roberta-wwm-ext` [Cui et al., 2020]. This tokenizer uses whole-word masking-aware subword vocabularies of size 21,128 and handles Chinese text directly without external word segmentation. All sequences are truncated or padded to a maximum length of 256 tokens, which is sufficient to cover the vast majority of validation examples without loss.

3 Model Architectures

3.1 Basic Model 1: Character-Level BiLSTM

The BiLSTM classifier treats each review as a sequence of character indices and processes it as follows.

1. **Embedding layer.** Each character index is mapped to a 64-dimensional trainable embedding vector.
2. **Encoder.** A single-layer bidirectional LSTM with hidden size 64 (total state dimension 128) encodes the sequence.
3. **Pooling.** The final forward and backward hidden states are concatenated into a 128-dimensional sentence vector.
4. **Regularisation.** Dropout with probability 0.3 is applied to the pooled vector.
5. **Classification head.** A linear layer maps the 128-dimensional vector to a 2-class logit.

The model has 260,546 total (and trainable) parameters. Its advantage is simplicity and fast training, but it cannot exploit pretrained linguistic knowledge and must learn all sentiment-relevant representations from supervised signal alone.

3.2 Basic Model 2: Baseline RoBERTa Classifier

The second model uses a pretrained Chinese RoBERTa encoder [Cui et al., 2020] as a fixed feature extractor with a simple trainable classification head.

1. **Encoder.** `hfl/chinese-roberta-wwm-ext`, a 12-layer Transformer with 768-dimensional hidden states and 12 attention heads.
2. **Pooling.** The [CLS] token representation from the final encoder layer is used as the sentence vector.
3. **Projection MLP.** A linear projection from 768 to a 512-dimensional hidden representation, followed by GELU activation and dropout with probability 0.08.
4. **Classification head.** A linear layer producing 2-class logits.

The encoder is fully frozen; only the projection MLP and classification head are trained, giving approximately 394,754 trainable parameters out of ~ 102 M total. Using a frozen encoder makes the model fast to train and avoids overfitting on the small task-specific dataset, but limits the backbone’s ability to adapt its representations to the sentiment domain.

3.3 Advanced Task: Baseline and Modified Architecture

Controlled frozen baseline. For the advanced task we construct a *matched frozen baseline* that shares the same backbone, optimiser schedule, hidden dimension (512), dropout (0.08), batch size (12), and random seed as the advanced model, but (i) freezes all encoder parameters and (ii) uses only the [CLS] token as the sentence representation. This baseline has 394,754 trainable parameters.

Advanced BERT architecture. The advanced classifier extends the baseline in five directions, each targeting a known limitation of the frozen [CLS]-head design.

1. **Layer mixing.** Rather than consuming only the final encoder layer, the model learns a softmax-normalised scalar mixture of the top 4 hidden layers. Concatenating representations across abstraction levels allows the classifier to leverage both fine-grained surface features (captured at lower layers) and coarser semantic compositionality (at higher layers) [Tenney et al., 2019].
2. **Multi-view pooling.** Four complementary sentence views are computed over the mixed representations:
 - [CLS] token: a global sequence summary distilled during pretraining;
 - masked mean pooling: stable average representation;
 - learned attention pooling: task-driven weighted sum;
 - masked max pooling: captures the most salient token per dimension, useful for polarity-anchoring content words.
3. **Gated fusion.** A learned gating network weights the four pooled views per example. The gated weighted sum is concatenated with the original views and layer-normalised before classification, giving the model input-adaptive sentence representation.
4. **Parameter-efficient adaptation.** LoRA adapters [Hu et al., 2022] are injected into the query, key, and value projections of all 12 encoder layers (rank 24, $\alpha = 32$), and into the attention output projection of the top 4 layers. BitFit [Ben Zaken et al., 2022] additionally enables all encoder bias parameters. LayerNorm parameters remain frozen. This allows the backbone to adapt to the sentiment domain while keeping the vast majority of parameters frozen.
5. **Wider classification head.** The fused representation is projected to a 512-dimensional hidden layer via GELU and dropout 0.08, increasing head capacity to match the richer input.

The advanced model has 105,730,571 total parameters but only 3,565,835 trainable parameters (3.37% of total), making it more parameter-efficient than the full fine-tuning alternative.

Table 2 summarises the architectural differences.

Table 2: Architecture comparison: controlled frozen baseline vs. advanced BERT.

Component	Frozen Baseline	Advanced BERT
Encoder usage	Final layer only, frozen	Learned mix of top-4 layers
Pooling	[CLS] only	[CLS] + mean + attn + max + gated fusion
Adaptation	None	LoRA (rank 24) + BitFit
Classifier head	512-dim MLP on [CLS]	512-dim MLP on fused representations
Trainable params	394,754	3,565,835

4 Training Setup

4.1 Common Implementation Details

All models are trained using the shared pipeline in `train.py` with PyTorch. The following choices apply uniformly: cross-entropy loss, AdamW optimiser [Loshchilov and Hutter, 2019], checkpoint selection by maximum validation macro F1, TensorBoard logging under `runs/`, and fixed random seed 42. The best checkpoint for each experiment is saved to `checkpoints/<experiment_name>/best_model.pt`.

4.2 Per-Model Hyperparameters

Table 3 lists the key hyperparameters for each experiment. Crucially, the controlled frozen baseline and the advanced BERT model share all non-architectural settings, so any observed difference in validation performance is attributable to the architectural changes described in Section 3.3.

Table 3: Training hyperparameters. The controlled baseline and advanced BERT share all settings except architectural ones (marked with †).

Hyperparameter	BiLSTM	Frozen Baseline	Advanced BERT
Optimizer	AdamW	AdamW	AdamW
Learning rate (head/main)	10^{-3}	5.13×10^{-5}	5.13×10^{-5}
Learning rate (encoder)	—	—	3.91×10^{-5}
Weight decay	0.05	0.00169	0.00169
Batch size	64	12	12
Epochs	15	4	4
Warmup ratio	—	0.150	0.150
Final LR scale	—	0.05	0.05
Dropout	0.30	0.08	0.08
Classifier hidden	—	512	512
Max sequence length	512 chars	256 tokens	256 tokens
LoRA rank / α †	—	—	24 / 32
Layer mix depth†	—	—	4
Encoder frozen†	—	Yes	No (LoRA+BitFit)

5 Experiments

5.1 Basic Task Comparison

The basic task pitches the character-level BiLSTM (Section 3.1) directly against the RoBERTa baseline (Section 3.2). The comparison is designed to measure the magnitude of the advantage conferred by large-scale language pretraining under the same dataset and evaluation conditions. Both models are trained to convergence according to their respective early-stopping criteria (validation macro F1 on held-out data).

5.2 Advanced Task Comparison

For the advanced task we follow a controlled experimental protocol. The frozen RoBERTa baseline is retrained with hyperparameters that exactly mirror those of the advanced model (see Table 3). The only intended differences are architectural: the matched baseline freezes the encoder and uses only the [CLS] token, while the advanced model adds layer mixing, multi-view pooling, gated fusion, and parameter-efficient encoder adaptation. This protocol does not individually ablate each advanced component, but it does cleanly isolate the aggregate effect of the full architectural upgrade from confounds related to optimiser schedule or regularisation.

6 Results

6.1 Main Quantitative Results

Table 4 reports the primary quantitative results. All metrics are measured on the held-out validation set using the best checkpoint of each run (selected by macro F1).

Table 4: Validation-set results. All values come directly from the saved best checkpoint of each run. Macro F1 is the primary selection metric.

Model	Acc.	Prec.	Recall	F1	Macro F1	Val. Loss
BiLSTM (char-level)	0.9252	0.9505	0.9575	0.9540	0.8770	0.2596
Frozen RoBERTa baseline	0.9172	0.9426	0.9560	0.9493	0.8622	0.2166
Advanced BERT	0.9580	0.9738	0.9743	0.9741	0.9317	0.1435

Table 4 reveals three clear findings. First, the BiLSTM achieves a respectable 92.52% accuracy and 0.8770 macro F1, demonstrating that character-level recurrent models can handle short Chinese reviews adequately. Second, the frozen RoBERTa baseline performs slightly below the BiLSTM on accuracy (91.72%) and macro F1 (0.8622), yet its validation loss is lower (0.2166 vs. 0.2596), suggesting better-calibrated predictions despite weaker raw classification. Third, the advanced BERT achieves the strongest results across all six metrics, with accuracy 4.08 points above the frozen baseline and macro F1 improved by 0.0695—a gap that far exceeds typical run-to-run variance.

6.2 Advanced Task: Gain over Frozen Baseline

Table 5 quantifies the per-metric improvement of the advanced model relative to the matched frozen baseline.

Table 5: Absolute metric improvement of Advanced BERT over the controlled frozen baseline (Δ = Advanced – Baseline; for loss, negative Δ is better).

	Acc.	Prec.	Recall	F1	Macro F1	Val. Loss
Δ	+0.0408	+0.0312	+0.0183	+0.0248	+0.0695	−0.0731

The largest relative improvement occurs on macro F1 (+0.0695), the metric most sensitive to minority-class behaviour. This suggests that the multi-view pooling and encoder adaptation primarily help the model better identify negative sentiment examples—the harder, minority class.

6.3 Training Curves and Learning-Rate Schedules

Figure 1 shows the training and validation curves for the controlled frozen baseline and the advanced BERT. The advanced model maintains a lower validation loss throughout training, and its validation macro F1 curve rises to a clearly higher plateau. The frozen baseline peaks at step 2502 (macro F1 = 0.8622), while the advanced model reaches its best checkpoint at step 2919 (macro F1 = 0.9317), indicating that the richer architecture continues to improve later into training.

Figure 2 shows the learning-rate schedules logged during training. Both models share the same warmup-plus-cosine-decay shape. The advanced model records a separate, lower LR for the encoder




Figure 1: Training dynamics for the controlled frozen baseline (blue) and the advanced BERT (orange). **Top left:** step-level training loss with raw signal (pale) and EMA-smoothed overlay (solid). **Top right:** validation loss at fixed evaluation intervals. **Bottom left:** validation macro F1. **Bottom right:** validation accuracy. Filled circles with annotations mark the best-checkpoint step for each run.

parameter group (green dashed), consistent with its intent to adapt the pretrained backbone gently without catastrophic forgetting.




Figure 2: Learning-rate schedules. **Orange solid:** adapter and classification-head LR for the advanced BERT. **Green dashed:** encoder LR for the advanced BERT, held lower to protect pretrained representations. The baseline main LR follows the same schedule as the adapter LR and is not drawn separately.

Figure 3 presents a compact bar-chart comparison of accuracy, macro F1, and validation loss at the best checkpoint of each model. The Δ annotations (green) confirm consistent improvement across all metrics, with the largest relative gain on macro F1 (+0.0695).




Figure 3: Best-checkpoint comparison of the three headline metrics for the controlled frozen baseline (blue) and the advanced BERT (orange). Higher is better for accuracy and macro F1; lower is better for validation loss. Green Δ labels quantify the direction and magnitude of improvement for the advanced model.

Figure 4 contextualises parameter efficiency across all three models. Despite sharing a ~ 102 M-parameter backbone, the advanced model activates only 3.57 M trainable parameters via LoRA and BitFit— $9\times$ more than the frozen baseline’s 395 K, yet still a small fraction of the cost of full fine-tuning.




Figure 4: Total (gray) vs. trainable (blue) parameter counts for each model. The BiLSTM is fully trainable but tiny (≈ 261 K). Both BERT-based models share a ~ 102 M-parameter backbone; the advanced model activates $9\times$ more trainable parameters than the frozen baseline by injecting LoRA adapters and enabling BitFit, while keeping the vast majority of the pretrained encoder frozen.

7 Discussion

7.1 Comparison of the Basic Models

The most striking result from the basic-task comparison is that the character-level BiLSTM *outperforms* the frozen RoBERTa baseline on accuracy and macro F1 despite having only 0.25% as many parameters. This outcome is not entirely surprising given the task context. The frozen baseline extracts only the [CLS] token and projects it to a decision boundary, but the classification head receives no gradient signal from the encoder—the pretrained representations are static. By contrast, the BiLSTM is fully trainable end-to-end, so every parameter is optimised directly for the binary sentiment objective.

We note, however, that the frozen baseline achieves lower validation loss (0.2166 vs. 0.2596), suggesting more calibrated predictions even when its discrete metrics are slightly lower. This

discrepancy is consistent with the hypothesis that the RoBERTa representations contain richer latent information than the frozen-head design can fully exploit.

Regarding convergence behaviour, the BiLSTM is trained for 15 epochs with a $1,000\times$ larger learning rate and converges stably within that budget, whereas the frozen baseline uses far fewer gradient updates (4 epochs) at a smaller learning rate. The architectures are therefore not trained under identical compute budgets, which should be kept in mind when comparing absolute metric values.

7.2 Effect of the Architectural Modification

Under the matched controlled protocol the advanced BERT improves macro F1 by 0.0695—a gain that is substantially larger than typical per-seed variance for this class of model. We attribute this improvement to three complementary mechanisms.

Layer mixing. Different encoder layers encode different linguistic abstractions [Tenney et al., 2019]. By taking a weighted mixture of the top 4 layers the advanced model can simultaneously use shallow syntactic signals and deep semantic compositionality, both of which are relevant for sentiment.

Multi-view pooling. A single pooling strategy inevitably discards information. Mean pooling smooths over polarity-bearing tokens; max pooling suppresses them in favour of a single extreme activating unit. By combining four complementary views—[CLS], mean, attention, and max—and learning to gate them per example, the model obtains a more complete sentence representation.

LoRA and BitFit adaptation. The pretrained encoder was optimised for a masked-language-modelling objective on a broad web corpus, not for binary sentiment classification on consumer reviews. LoRA and BitFit provide a parameter-efficient channel through which the encoder can adjust to the target distribution without catastrophically forgetting its pretrained representations. The improvement in minority-class recall (negative sentiment) is consistent with this interpretation: domain adaptation helps the model activate features specific to the less frequent class.

It is important to note that the advanced architecture bundles several modifications together, so we cannot attribute the full +0.0695 gain to any single component. A full per-component ablation remains future work.

7.3 Limitations

Several methodological limitations affect the strength of the conclusions. First, all experiments are single-seed; variance across random seeds is unknown. Second, the class imbalance (81% positive) is not addressed by any explicit balancing strategy such as class-weighted loss or oversampling, which could further improve minority-class performance. Third, while the advanced task comparison is carefully controlled at the optimiser level, the advanced model was initially designed with the benefit of a preceding hyperparameter search, which may introduce implicit selection bias. Fourth, the evaluation is restricted to one dataset, limiting the generalisability of the conclusions.

8 Conclusion

We presented an empirical evaluation of three sentiment classifiers on a Chinese binary review dataset. The BiLSTM provides a competitive character-level baseline with strong macro F1 (0.8770) under fully supervised training. The frozen RoBERTa baseline demonstrates that pretrained representations alone are not sufficient without end-to-end adaptation. Most importantly, the advanced BERT classifier—which combines layer mixing, gated multi-view pooling, and parameter-efficient LoRA and BitFit adaptation—achieves validation macro F1 of 0.9317 and accuracy of 95.80%, substantially outperforming the matched frozen baseline under identical non-architectural conditions.

The key lesson is that *how* encoder features are aggregated and *where* adaptation is permitted inside a pretrained network both matter substantially. A naive frozen-head design leaves significant performance on the table, even when the backbone carries rich linguistic knowledge.

Future work should systematically ablate each advanced component to quantify individual contributions; evaluate robustness across multiple random seeds; explore class-balanced training to further reduce minority-class error; and test transferability of the advanced architecture to other Chinese sentiment or topic classification datasets.

References

- Elad Ben Zaken, Yoav Goldberg, and Shauli Ravfogel. BitFit: Simple parameter-efficient fine-tuning for transformer-based masked language-models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 1–9, 2022. doi: 10.18653/v1/2022.acl-short.1. URL <https://aclanthology.org/2022.acl-short.1/>.
- Yiming Cui, Wanxiang Che, Ting Liu, Bing Qin, Shijin Wang, and Guoping Hu. Revisiting pre-trained models for Chinese natural language processing. *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 657–668, 2020. doi: 10.18653/v1/2020.findings-emnlp.58. URL <https://aclanthology.org/2020.findings-emnlp.58/>.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186. Association for Computational Linguistics, 2019. doi: 10.18653/v1/N19-1423. URL <https://aclanthology.org/N19-1423/>.
- Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8): 1735–1780, 1997.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.
- Ian Tenney, Dipanjan Das, and Ellie Pavlick. BERT rediscovers the classical NLP pipeline. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4593–4601, 2019. doi: 10.18653/v1/P19-1452. URL <https://aclanthology.org/P19-1452/>.

A Checkpoint Metadata

The following checkpoint directories contain the model weights used to compute all reported metrics.

- checkpoints/lstm_classifier/best_model.pt
- checkpoints/bert_classifier/best_model.pt
- checkpoints/bert_classifier_advanced/best_model.pt

TensorBoard logs are stored under runs/ and can be visualised with:

```
tensorboard --logdir ./runs
```

The most relevant event files are

- runs/bert_classifier/bert_classifier_20260331-220647 (controlled frozen baseline, best at step 2502, macro F1 = 0.8622)
- runs/bert_classifier_advanced/bert_classifier_advanced_20260331-025421 (advanced BERT reproduced run, best at step 2919, macro F1 = 0.9317)

B Implementation Notes

All experiments were run with PyTorch via the uv virtual environment pinned to Python 3.11. Figure generation is handled by `scripts/export_report_figures.py`, which reads TensorBoard event files and writes publication-quality .pdf and .png exports under figures/.